

"Introduction to Machine Learning" Exam (November 2023)

Guidelines:

- In the *Answer Sheet*, circle the letter (only one) corresponding to the correct answer.
- Each correct answer **adds one point**. Each incorrect answer **subtracts half a point**.


1. The height of a person is:
 - a. An observable feature.
 - b. A latent feature.
 - c. A handcrafted feature.
2. Which of the following statements about **representation learning** is incorrect?
 - a. It is computationally convenient at test time.
 - b. It aims at learning latent features.
 - c. It aims at learning handcrafted features.
3. Which of the following loss functions is not adapted to a classification task?
 - a. $\ell(h, z = (x, y)) = e^{-yh(x)}$.
 - b. $\ell(h, z = (x, y)) = (y - h(x))^2$.
 - c. $\ell(h, z = (x, y)) = [yh(x) < 0]$, where $[\cdot]$ is an indicator function.
4. Which loss function can be subject to gradient explosion?
 - a. Hinge loss.
 - b. Exponential loss.
 - c. (0-1)-loss.
5. A classifier h_θ is obtained from S by solving the following optimization problem:
$$\min_{\theta} \sum_{i=1}^n \ell(h_\theta, S) + \lambda \|\theta\|_p$$

h_θ has a training error=2% and a test error=15%. What do you suggest?

 - a. Decrease the number of training examples n ?
 - b. Use another loss function ℓ .
 - c. Increase the value of the regularization parameter λ .
6. Which of the following statements about the regularization parameter λ is correct?
 - a. A too large λ can cause your hypothesis to underfit the data.
 - b. A too large λ can cause your hypothesis to overfit the data.
 - c. Using a large value of λ cannot hurt the performance of your hypothesis.

7. Which ℓ_p -norm is convex?
- a. ℓ_0
 - b. $\ell_{0.5}$
 - c. ℓ_1 .
8. The use of the ℓ_1 -norm as a regularization term allows to:
- a. get a sparse model only.
 - b. reduce the risk of overfitting only.
 - c. both get a sparse model and reduce the risk of overfitting.
9. In a binary classification setting, an error occurs when:
- a. $yh(x) \neq 0$.
 - b. $y - h(x) < 0$.
 - c. $yh(x) < 0$.
10. Given the objective function $J(\theta) = \theta^2$. Given an initialization $\theta = -1$ and a learning rate $\alpha = 0.3$, what is the updated value of θ after the first iteration of gradient descent?
- a. $\theta = -1.6$.
 - b. $\theta = -0.4$.
 - c. $\theta = 0.8$.
11. As the number of features grows, the amount of data to learn well grows:
- a. linearly fast.
 - b. quadratically fast.
 - c. exponentially fast.

12. In the following figure, which hypothesis would be selected according to Occam's razor principle?



- a. h_1 .
 - b. h_2 .
 - c. h_3 .
13. What does the term PAC mean in statistical learning theory?
- a. Probably Almost Correct.
 - b. Probably Approximately Correct.
 - c. Potentially Approximately Correct.
14. The optimal classifier h^* is usually not selected by the learning algorithm because:
- a. It does not depend on the training examples.
 - b. It changes according to the algorithm.
 - c. It is not in the family of hypotheses.
15. Let $f_A(x) = x, \forall x \in [0, \sqrt{2}]$ and $f_B(x) = 1, \forall x \in [0, 1]$ be the distributions of the class A and B , respectively. What is the value of the bayesian error ϵ_B ?
- a. $\epsilon_B = \frac{1}{4}$
 - b. $\epsilon_B = \frac{1}{2}$
 - c. $\epsilon_B = \frac{2}{3}$
16. A training set S of 4 points is shattered by a family of hypotheses \mathcal{H} if \mathcal{H} realizes:
- a. the 4 dichotomies of S .
 - b. the 8 dichotomies of S .
 - c. the 16 dichotomies of S .
17. Which of the following criteria is not a complexity measure?
- a. The VC-dimension.
 - b. The number of training examples.
 - c. The Rademacher complexity.

18. The error obtained using K -cross-validation is an estimate of the generalization error. Is this estimate unbiased?
- Yes .
 - No
 - It depends on K
19. Regarding bias and variance, which of the following statements is right?
- Models which overfit have a high bias.
 - Models which overfit have a low bias. .
 - Models which underfit have a high variance.
20. What is the Hessian matrix H of the function $f(x, y) = -xy + 2x^2y^4 - 3y^3$?
- $H = \begin{pmatrix} 4x^2 & 16xy^3 \\ 16y^3 & 12x^2y^2 - y \end{pmatrix}$
 - $H = \begin{pmatrix} 4x^2 & 10xy^3 - 1 \\ 10xy^3 - 1 & 12x^2y^2 - y \end{pmatrix}$
 - $H = \begin{pmatrix} 4y^3 & 16xy^3 - 1 \\ 16xy^3 - 1 & 24x^2y^2 - 18y \end{pmatrix}$
21. What is the incorrect statement?
- a k -NN approximates the bayesian classifier
 - a k -NN learns spheres. .
 - a k -NN is lazy.
22. With noisy data, which of the following options would you consider for a k -NN classifier?
- Increase the value of k .
 - Decrease the value of k
 - Dealing with noise is independent from the value of k
23. What happens when running a k -NN algorithm on imbalanced datasets composed of a few positives and many negatives?
- The decision boundaries are overly in favor of the positive class.
 - The decision boundaries are not affected by the imbalance.
 - The decision boundaries are overly in favor of the negative class.

24. With a kNN classifier, under which conditions $\frac{k_n}{nV_n}$ is a good estimate of $p(\mathbf{x})$, where V_n is a volume centered at \mathbf{x} , k_n is the number of examples in V_n and n is the size of the training set?

- a. $\lim_{n \rightarrow \infty} V_n = 0$, $\lim_{n \rightarrow \infty} k_n = \infty$ and $\lim_{n \rightarrow \infty} \frac{k_n}{n} = 0$
- b. $\lim_{n \rightarrow \infty} V_n = 0$, $\lim_{n \rightarrow \infty} k_n = 0$ and $\lim_{n \rightarrow \infty} \frac{k_n}{n} = 0$
- c. $\lim_{n \rightarrow \infty} V_n = 0$, $\lim_{n \rightarrow \infty} k_n = 0$ and $\lim_{n \rightarrow \infty} \frac{k_n}{n} = \infty$

25. Which statement is wrong?

- a. $\epsilon_{1NN} \leq \epsilon_B$
- b. $\epsilon_B \leq 2\epsilon_{1NN}$
- c. $\epsilon_{1NN} \leq 2\epsilon_B$

26. A training set is composed of three examples $(x_1, x_2) \in \mathbb{R}^2$, where the label $y \in \{-1, +1\}$.

x_1	x_2	y
a	c	+1
b	d	+1
a	b	-1

Using the naive Bayes classifier, what is the value of $P(y = +1 | x_1 = a, x_2 = d)$?

- a. 0.
- b. $\frac{1}{3}$.
- c. $\frac{1}{2}$.


27. Which of the following statements about the naive Bayes classifier is wrong?

- a. It overcomes the problem of curse of dimensionality.
- b. It assumes that the features are not independent.
- c. It allows to smooth the probabilities.

28. The Bayesian error is supposed to be $\epsilon_B = 0\%$. A classifier h has a training error=9% and a validation error=10%. What do you suggest?

- a. To increase the number of training examples.
- b. To increase the complexity of your classifier.
- c. To reduce the complexity of your classifier.

29. Given the following ranking obtained from 3 positive examples (in red) and 5 negatives (in blue). What is the value of the Average Precision AP?



- a. $\frac{5}{9}$.
- b. $\frac{2}{11}$.
- c. $\frac{2}{7}$.
30. Which of the following is incorrect about the F_β -measure?
- a. The F_β -measure is not convex.
- b. The F_β -measure is defined as follows: $F_\beta = \frac{(1+\beta^2)Precision+Recall}{\beta^2 \times Precision \times Recall}$
- c. F_β is a non separable function.